

Homework 4: Lazy Memory Allocation

As we discussed in class, one of the many cool *tricks* that paging hardware can be used for is *lazy allocation* -- only giving out physical memory when it's used, rather than when it's allocated. In this assignment, we'll modify xv6 to use a lazy allocator. We won't cover every corner case, but by the end of it we will have a basic working implementation.

Getting the Code from Github

As with last time, we'll be working off of a slightly modified version of xv6. The major difference is that the `cat` program has been modified so that it uses dynamic memory allocation rather than static allocation, which will allow us to uncover an edge case in a naive implementation.

If you still have your xv6 directory from last time, remove or rename it first. Then get the base xv6 code for this assignment:

```
$ git clone https://github.com/moyix/xv6-public.git
Cloning into 'xv6-public'...
remote: Counting objects: 4501, done.
remote: Compressing objects: 100% (17/17), done.
remote: Total 4501 (delta 3), reused 0 (delta 0), pack-reused 4484
Receiving objects: 100% (4501/4501), 11.68 MiB | 3.46 MiB/s, done.
Resolving deltas: 100% (1800/1800), done.
Checking connectivity... done.
$ cd xv6-public/
$ git checkout hw5
Branch hw5 set up to track remote branch hw5 from origin.
Switched to a new branch 'hw5'
```

Part 1: Use the debugger

In order to get acquainted with the debugger, we will start with debugging XV6. XV6 has a Makefile rule to make debugging simple, however you will need to have two open command line windows. Type the following commands on each:

- Command Line 1: `$ make qemu-gdb`
- Command Line 2: `$ gdbtui kernel`

Command line 1 is going to basically create a debuggeable version of XV6. Command line 2, is going to start the debugger.

Now we will setup a breakpoint in the system call that allocates memory. The call that user-process use in xv6 to allocate memory is called `sbrk()`; it's kernel-mode implementation is `sys_sbrk()` in `sysproc.c`. **Set a breakpoint in this function.** In order to do that, you might need a gdb cheatsheet like the one in NYU classes. Then answer the following questions when you submit your patch:

1. Run a user process like: 'ls'
2. What is the size of `n` when your breakpoint gets hit?
3. Print out the stackframe (backtrace) for this call.

Part 2: Removing Eager Allocation

Change `sys_sbrk()` so that it just adjusts `proc->sz` and returns the old `proc->sz` (i.e., remove the call to `growproc()`).

After rebuilding xv6, try running a command like `echo hello`:

```
init: starting sh
$ echo hello
pid 3 sh: trap 14 err 6 on cpu 0 eip 0x12bb addr 0x4004--kill proc
```

What went wrong, and why? (Answering this question is not part of the assignment, just think about it).

Part 3: Implementing Lazy Allocation

The message above comes from the `trap()` function in `trap.c`. It indicates that an exception (number 14, which corresponds to a page fault on x86) was raised by the CPU; it happened when `eip` (the program counter) was `0x12bb`, and the faulting address was `0x4004`. It then kills the process by setting `proc->killed` to 1. Can you find in the code in `trap.c` where that happened? (Answering this question is not part of the assignment either, just getting you familiar with the code)

Add code to `trap()` to recognize this particular exception and allocate and map the page on-demand.

Once you have implemented the allocation, try running `echo hello` again. It should work normally.

Hints:

1. The constant `T_PGFLT`, defined in `traps.h`, corresponds to the exception number for a page fault.
2. The virtual address of the address that triggered the fault is available in the `cr2` register; xv6 provides the `pcr2()` function to read its value.
3. Look at the code in `allocuvn()` to see how to allocate a page of memory and map it to a specific user address.
4. Remember that the first access might be in the middle of the page, and so you'll have to round down to the nearest `PGSIZE` bytes. xv6 has a handy function called `PGROUNDDOWN` you can use for this.
5. You will need to use the `mappages()` function, which is declared as `static`, meaning it can't be seen from other C files. You'll need to make it `non-static`, and then add its `prototype` to some header file that is included by both `trap.c` and `vm.c` (`defs.h` is a good choice).

Part 4: Handling Some Edge Cases

Although simple commands like `echo` work now, there are still some things that are broken. For example, try running `cat README`. Depending on how you implemented Part 2, you may see:

```
init: starting sh
$ cat README
unexpected trap 14 from cpu 1 eip 80105282 (cr2=0x3000)
cpu1: panic: trap
      8010683a 801064fa 80101f4e 80101174 80105725 8010556a 801066f9 801064fa 0 0
```

Why is this happening? Debug the problem and find a fix (if this part works already, you don't need to make any changes).

Next, let's see if pipes work, by running `cat README | grep the`

```
$ cat README | grep the
cpu0: panic: copyvm: page not present
      8010828e 80104693 8010630b 8010556a 801066f9 801064fa 0 0 0 0
```

Find out what's going on, and implement a fix.

Hints

1. Think about what happens if the kernel tries to read an address that hasn't been mapped yet (for example, if we use `sbrk()` to allocate some memory and then hand that buffer to the `read()` syscall). Read the code in `trap()` with that case in mind.
2. Making pipes work is a very tiny change -- just one line in one file. Don't overthink it.

Submitting

As with HW4, you will use git to create a patch.

Commit your changes:

```
$ git commit --all --message="Implement lazy allocation"
[hw5 ef751c0] Implement lazy allocation
4 files changed, 30 insertions(+), 10 deletions(-)
```

(Note: if you added any new files, you will also have to use `git add <filename>` before you run `git commit`.)

Now create the patch file:

```
$ git format-patch hw5.unmodified
0001-Implement-lazy-allocation.patch
```

The command creates a file, `0001-Implement-lazy-allocation.patch`, containing the changes you've made. Submit this file on NYU Classes

Final Notes

Although this covers most use cases, there are still a few stray things that it doesn't handle:

1. Negative arguments to `sbrk()` should reduce the size of the process and deallocate the pages.
2. This code will allow multiple processes to collectively allocate more memory than the system has (*overcommit*).

It may be helpful, if you want to understand the material better, to think about how you could fix these issues.

Credit: Adapted from a homework by Brendan-Dolan Gavit.

Credit: Adapted from MIT's 6.828 in-class exercise [xv6 lazy page allocation](#)